

```

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    print(os.path.join(dirname))

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,Dropout,Convolution2D,MaxPooling2D,Flatten #action detection
import tensorflow as tf
import matplotlib.pyplot as plt
from IPython.display import HTML

from tensorflow.keras.preprocessing.image import ImageDataGenerator
IMAGE_SIZE = 128

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    horizontal_flip=True
)
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Cars Dataset/train',
    target_size=(IMAGE_SIZE,IMAGE_SIZE),
    class_mode="sparse",
)

Found 3352 images belonging to 7 classes.

count=0
for image_batch, label_batch in train_generator:
    # print(label_batch)
    print(image_batch[0])
    break

[[[0.22543551 0.22543551 0.17837669]
 [0.21077232 0.21077232 0.16371349]
 [0.4016731 0.4016731 0.3546143 ]
 ...
 [0.9568628 0.9568628 0.91300875]
 [0.9584651 0.9584651 0.91604483]
 [0.9607844 0.9607844 0.91372555]]

[[[0.3520475 0.3520475 0.30498868]
 [0.2639154 0.2639154 0.21685658]
 [0.3979375 0.3979375 0.3508787 ]
 ...
 [0.9568628 0.9568628 0.91295815]
 [0.95843977 0.95843977 0.9160701 ]
 [0.9607844 0.9607844 0.91372555]]

[[[0.44917092 0.44917092 0.40211207]
 [0.30422154 0.30422154 0.2571627 ]
 [0.39525118 0.39525118 0.34819236]
 ...
 [0.9568628 0.9568628 0.9129075 ]
 [0.9584145 0.9584145 0.91609544]
 [0.9607844 0.9607844 0.91372555]]

...

[[[0.09320191 0.08928034 0.11280976]
 [0.09327786 0.08935629 0.11288571]
 [0.0933538 0.08943223 0.11296164]
 ...
 [0.21592927 0.2120077 0.19632143]
 [0.21854743 0.21462587 0.19893959]
 [0.2182389 0.21431734 0.19863106]]

[[[0.09573017 0.0918086 0.11533801]
 [0.09570486 0.09178329 0.1153127 ]
 [0.09567954 0.09175797 0.11528739]
 ...

```

```

[0.20134674 0.19742517 0.18173888]
[0.20417634 0.20025477 0.1845685 ]
[0.20599721 0.20207565 0.18638937]]

[[0.08026382 0.07634225 0.09987167]
 [0.08011194 0.07619037 0.09971979]
 [0.07996006 0.07603849 0.09956791]
 ...
 [0.22976013 0.22583856 0.21015228]
 [0.23368162 0.22976005 0.21407378]
 [0.23067619 0.22675462 0.21106835]]]

class_names = list(train_generator.class_indices.keys())
class_names

['Audi',
 'Hyundai Creta',
 'Mahindra Scorpio',
 'Rolls Royce',
 'Swift',
 'Tata Safari',
 'Toyota Innova']

test_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    horizontal_flip=True)

test_generator = test_datagen.flow_from_directory(
    '/content/drive/MyDrive/Cars Dataset/test',
    target_size=(IMAGE_SIZE, IMAGE_SIZE),
    class_mode="sparse"
)

Found 813 images belonging to 7 classes.

for image_batch, label_batch in test_generator:
    print(image_batch[0])
    break

```

```

[[[1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 ...
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]]]

[[[1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 ...
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]]]

[[[1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 ...
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]]]

...

[[0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 ...
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]]]

[[0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 ...
 [1.      1.      1.      ]
 [1.      1.      1.      ]
 [1.      1.      1.      ]]]

[[0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 [0.9960785 0.9960785 0.9960785]
 ...

```

```
[1.      1.      1.      ]
[1.      1.      1.      ]
[1.      1.      1.      ]]
```

```
sz = 128
```

```
# Initializing the CNN
```

```
model = Sequential()
```

```
# First convolution layer and pooling
```

```
model.add(Convolution2D(32, (3, 3), input_shape=(sz, sz, 3), activation='relu'))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
# Second convolution layer and pooling
```

```
model.add(Convolution2D(32, (3, 3), activation='relu'))
```

```
# input_shape is going to be the pooled feature maps from the previous convolution layer
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
# Flattening the layers
```

```
model.add(Flatten())
```

```
# Adding a fully connected layer
```

```
model.add(Dense(units=97, activation='relu'))
```

```
model.add(Dropout(0.40))
```

```
model.add(Dense(units=32, activation='relu'))
```

```
model.add(Dense(units=7, activation='softmax')) # softmax for more than 2
```

```
model.summary()
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_9 (Conv2D)	(None, 61, 61, 32)	9248
max_pooling2d_9 (MaxPooling2D)	(None, 30, 30, 32)	0
flatten_4 (Flatten)	(None, 28800)	0
dense_12 (Dense)	(None, 97)	2793697
dropout_4 (Dropout)	(None, 97)	0
dense_13 (Dense)	(None, 32)	3136
dense_14 (Dense)	(None, 7)	231
=====		
Total params: 2807208 (10.71 MB)		
Trainable params: 2807208 (10.71 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
model.compile(optimizer='adam', loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False), metrics=['accuracy'])
```

```
history = model.fit(
    train_generator,
    validation_data=test_generator,
    epochs=5
)
```

```
Epoch 1/5
105/105 [=====] - 314s 3s/step - loss: 1.8184 - accuracy: 0.2915 - val_loss: 1.7538 - val_accuracy: 0.3764
Epoch 2/5
105/105 [=====] - 84s 797ms/step - loss: 1.6880 - accuracy: 0.3729 - val_loss: 1.5904 - val_accuracy: 0.46
Epoch 3/5
105/105 [=====] - 87s 822ms/step - loss: 1.5617 - accuracy: 0.4335 - val_loss: 1.4325 - val_accuracy: 0.50
Epoch 4/5
105/105 [=====] - 87s 822ms/step - loss: 1.4318 - accuracy: 0.4717 - val_loss: 1.3584 - val_accuracy: 0.51
Epoch 5/5
105/105 [=====] - 85s 809ms/step - loss: 1.3464 - accuracy: 0.5075 - val_loss: 1.3073 - val_accuracy: 0.53
```

```
scores = model.evaluate(test_generator)
```

```
26/26 [=====] - 10s 382ms/step - loss: 1.3057 - accuracy: 0.5301
```

scores

```
[1.3057115077972412, 0.5301352739334106]
```

```
history.history.keys()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
type(history.history['loss'])
```

```
list
```

```
len(history.history['loss'])
```

```
5
```

```
history.history['loss'][:5] # show loss for first 5 epochs
```

```
[1.8184212446212769,  
1.688022494316101,  
1.5616528987884521,  
1.4318387508392334,  
1.3464020490646362]
```

```
acc = history.history['accuracy']
```

```
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']
```

```
val_loss = history.history['val_loss']
```

```
import matplotlib.pyplot as plt
```

```
EPOCHS = 5
```

```
plt.figure(figsize=(5, 5))
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(range(EPOCHS), acc, label='Training Accuracy')
```

```
plt.plot(range(EPOCHS), val_acc, label='Validation Accuracy')
```

```
plt.legend(loc='lower right')
```

```
plt.title('Training and Validation Accuracy')
```

```
plt.subplot(1, 2, 2)
```


```
plt.plot(range(EPOCHS), loss, label='Training Loss')
```

```
plt.plot(range(EPOCHS), val_loss, label='Validation Loss')
```

```
plt.legend(loc='upper right')
```

```
plt.title('Training and Validation Loss')
```

```
plt.show()
```



```
def predict(model, img):
```

```
    img_array = tf.keras.preprocessing.image.img_to_array(images[i])
```

```
    img_array = tf.expand_dims(img_array, 0)
```

```
    predictions = model.predict(img_array)
```

```
    predicted_class = class_names[np.argmax(predictions[0])]
```

```
confidence = round(100 * (np.max(predictions[0])), 2)
return predicted_class, confidence
```

```
plt.figure(figsize=(15, 15))
for images, labels in test_generator:
    for i in range(6):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i])

        predicted_class, confidence = predict(model, images[i])
        actual_class = class_names[int(labels[i])]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confidence: {confidence}%")

    plt.axis("off")
break

1/1 [=====] - 0s 116ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 28ms/step
```

Actual: Tata Safari,
Predicted: Toyota Innova.
Confidence: 26.96%



Actual: Tata Safari,
Predicted: Tata Safari.
Confidence: 33.78%




Actual: Rolls Royce,
Predicted: Audi.
Confidence: 50.59%



Actual: Hyundai Creta,
Predicted: Toyota Innova.
Confidence: 57.19%



Actual:
Predicted:
Confidence



Actual:
Predicted:
Confidence

