

```

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
"""
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

"""

'\nimport os\nfor dirname, _, filenames in os.walk('/kaggle/input'):\n    for filename in filenames:\n        print(os.path.join(dirname, filename))\n\n'

import numpy as np
import pandas as pd
import os
from pathlib import Path

# Image processing
import cv2

# Utility
from tqdm import tqdm

# Sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

# Charts
import matplotlib.pyplot as plt
import seaborn as sns

# Tensorflow
from keras.preprocessing.image import ImageDataGenerator
from keras.applications import vgg16
from keras.layers import Dense, Dropout, BatchNormalization, Flatten
from keras.models import Model, load_model
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint, EarlyStopping

input_path = '../content/drive/MyDrive/Fish_Dataset'
# Create Series contains all file_paths
file_paths = pd.Series(list(Path(input_path).glob('**/*.png')), name='file_path')
# Convert posixpath to str
file_paths = file_paths.map(lambda file_path: str(file_path))
# Create Series contains all labels corresponding to file_path
labels = pd.Series(list(map(lambda file_path: file_path.split('/')[6], file_paths)), name='category')
# Create DataFrame contains file_path and label
fishes = pd.concat([file_paths, labels], axis=1)
# Remove all files in GT folder
fishes = fishes[fishes.category.map(lambda cate: cate[-2:] != 'GT')]
fishes = fishes.sample(frac=1).reset_index(drop=True)
fishes.head()

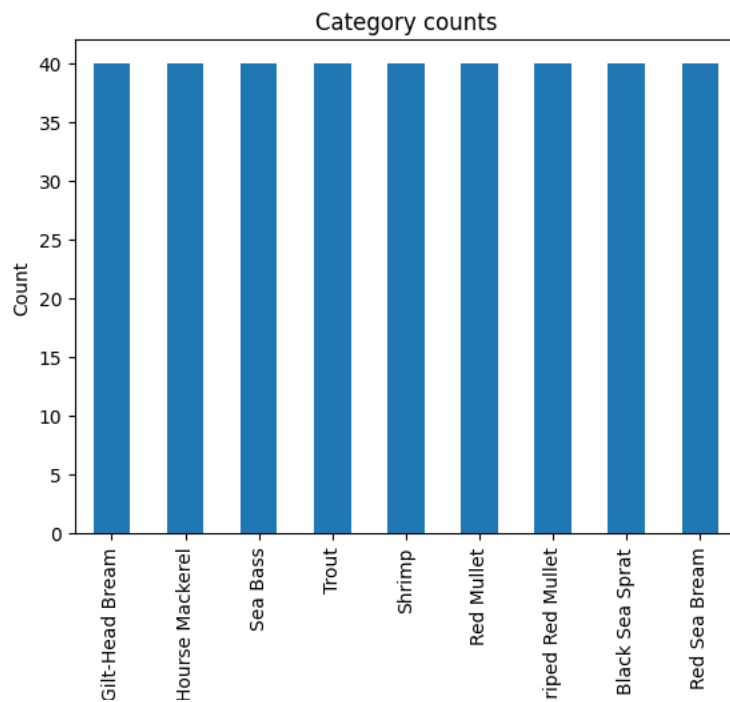

```

	file_path	category
0	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Gilt-Head Bream
1	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Horse Mackerel
2	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass
3	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass
4	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Trout

```


fishes.category.value_counts().plot(kind='bar')
plt.title('Category counts')
plt.xlabel('Category')
plt.ylabel('Count')
plt.show()

```



```
def plot_first_six_image():
    fig = plt.figure(figsize=(20, 20))
    for i in range(6):
        row = fishes.loc[i]
        ax = fig.add_subplot(1, 6, i+1)
        img = plt.imread(row.file_path)
        img = cv2.resize(img, (224, 224))
        ax.set_title(row.category)
        ax.imshow(img)
```

```
plot_first_six_image()
```



```
# Split to train/test dataset
train_fishes, test_fishes = train_test_split(fishes, test_size=0.2)
print(train_fishes)
print(test_fishes)
```

5

	file_path	category
188	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Trout
106	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass
352	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Striped Red Mullet
8	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Shrimp
100	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Trout
..
247	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Black Sea Sprat
267	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Red Sea Bream
309	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Black Sea Sprat
291	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Red Sea Bream
83	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass

[288 rows x 2 columns]

	file_path	category
51	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Gilt-Head Bream
42	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Shrimp
324	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Red Mullet
32	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Red Mullet
61	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass
..
184	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Black Sea Sprat
284	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Red Sea Bream
20	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Shrimp
186	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Trout
328	../content/drive/MyDrive/Fish_Dataset/Fish_Dat...	Sea Bass

[72 rows x 2 columns]

```

train_gen = ImageDataGenerator(validation_split=0.2)
test_gen = ImageDataGenerator()
train_batches = train_gen.flow_from_dataframe(dataframe=train_fishes,
                                              x_col='file_path',
                                              y_col='category',
                                              class_mode='categorical',
                                              batch_size=64,
                                              target_size=(224, 224),
                                              subset='training',
                                              color_mode='rgb')

valid_batches = train_gen.flow_from_dataframe(dataframe=train_fishes,
                                              x_col='file_path',
                                              y_col='category',
                                              class_mode='categorical',
                                              batch_size=64,
                                              target_size=(224, 224),
                                              subset='validation',
                                              color_mode='rgb')

# In test_batches, we should set shuffle=False to match predicted outputs and ground-truth labels later
test_batches = test_gen.flow_from_dataframe(dataframe=test_fishes, x_col='file_path',
                                           y_col='category',
                                           class_mode='categorical',
                                           batch_size=64,
                                           target_size=(224, 224),
                                           color_mode='rgb',
                                           shuffle=False)

```

Found 231 validated image filenames belonging to 9 classes.
 Found 57 validated image filenames belonging to 9 classes.
 Found 72 validated image filenames belonging to 9 classes.

```

base_model = vgg16.VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
if os.path.exists('../models') == False:
    os.mkdir('../models')
base_model.save('../models/vgg16_as_feature_extractor.h5')

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_n58889256/58889256 [=====] - 0s 0us/step
 /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3000: UserWarning: You are saving your model as an HDF5 file v saving_api.save_model(
 WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty

```

base_model = load_model('../models/vgg16_as_feature_extractor.h5')
for layer in base_model.layers:
    layer.trainable = False
base_model.summary()

```

WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.
 Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

```

block5_conv3 (Conv2D)      (None, 14, 14, 512)      2359808
block5_pool (MaxPooling2D) (None, 7, 7, 512)        0

```

```

=====
Total params: 14714688 (56.13 MB)
Trainable params: 0 (0.00 Byte)
Non-trainable params: 14714688 (56.13 MB)

```

```

last_layer = base_model.get_layer('block5_pool')
last_output = last_layer.output

```

```

x = Flatten()(last_output)
x = Dense(64, activation='relu', name='FC_1')(x)
x = BatchNormalization()(x)
x = Dropout(0.3)(x)
x = Dense(9, activation='softmax', name='softmax')(x)

```

```

new_model = Model(inputs=base_model.input, outputs=x)
new_model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
FC_1 (Dense)	(None, 64)	1605696
batch_normalization (Batch Normalization)	(None, 64)	256
dropout (Dropout)	(None, 64)	0
softmax (Dense)	(None, 9)	585

```

=====
Total params: 16321225 (62.26 MB)
Trainable params: 1606409 (6.13 MB)
Non-trainable params: 14714816 (56.13 MB)

```

```

early_stopping = EarlyStopping(monitor='val_accuracy',
                                min_delta=0,
                                patience=3,
                                verbose=1,
                                mode='auto',

```

```

        restore_best_weights=True)
new_model.compile(Adam(lr=0.0001), loss='categorical_crossentropy', metrics=['accuracy'])
history = new_model.fit_generator(generator=train_batches,
                                steps_per_epoch=train_batches.n // train_batches.batch_size,
                                validation_data=valid_batches,
                                validation_steps=valid_batches.n // valid_batches.batch_size,
                                epochs=10,
                                verbose=1,
                                callbacks=[early_stopping])
new_model.save('../models/fish_classifier_best_model.h5')

WARNING:absl:lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g., tf.keras.optimizer
<ipython-input-13-969f90e2b095>:8: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please
    history = new_model.fit_generator(generator=train_batches,
Epoch 1/10
3/3 [=====] - ETA: 0s - loss: 2.3713 - accuracy: 0.3174 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 129s 43s/step - loss: 2.3713 - accuracy: 0.3174
Epoch 2/10
3/3 [=====] - ETA: 0s - loss: 0.8071 - accuracy: 0.7605 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 105s 39s/step - loss: 0.8071 - accuracy: 0.7605
Epoch 3/10
3/3 [=====] - ETA: 0s - loss: 0.6214 - accuracy: 0.8563 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 107s 33s/step - loss: 0.6214 - accuracy: 0.8563
Epoch 4/10
3/3 [=====] - ETA: 0s - loss: 0.4630 - accuracy: 0.9042 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 105s 41s/step - loss: 0.4630 - accuracy: 0.9042
Epoch 5/10
3/3 [=====] - ETA: 0s - loss: 0.3424 - accuracy: 0.9102 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 105s 41s/step - loss: 0.3424 - accuracy: 0.9102
Epoch 6/10
3/3 [=====] - ETA: 0s - loss: 0.2870 - accuracy: 0.9688 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 122s 40s/step - loss: 0.2870 - accuracy: 0.9688
Epoch 7/10
3/3 [=====] - ETA: 0s - loss: 0.2218 - accuracy: 0.9701 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 105s 32s/step - loss: 0.2218 - accuracy: 0.9701
Epoch 8/10
3/3 [=====] - ETA: 0s - loss: 0.1793 - accuracy: 0.9880 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 108s 41s/step - loss: 0.1793 - accuracy: 0.9880
Epoch 9/10
3/3 [=====] - ETA: 0s - loss: 0.1747 - accuracy: 0.9948 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 121s 39s/step - loss: 0.1747 - accuracy: 0.9948
Epoch 10/10
3/3 [=====] - ETA: 0s - loss: 0.1547 - accuracy: 0.9896 WARNING:tensorflow:Early stopping conditioned on metric `loss` which has not been
3/3 [=====] - 123s 41s/step - loss: 0.1547 - accuracy: 0.9896

```

```

def plot_learning_curves(value='loss'):
    plt.plot(history.history[value], label='train')
    plt.plot(history.history['val_' + value], label='val')
    plt.xlabel('epochs')
    plt.ylabel(value + 'val_')
    plt.legend()
    plt.show(plot_learning_curves)

y_hat = new_model.predict(test_batches, verbose=1)
# Compare first 10 predicted value and ground truth label
print('First 10 predicted value from model: ', np.argmax(y_hat[:10], axis=1))
print('First 10 ground truth label: ', test_batches.labels[:10])

```

```

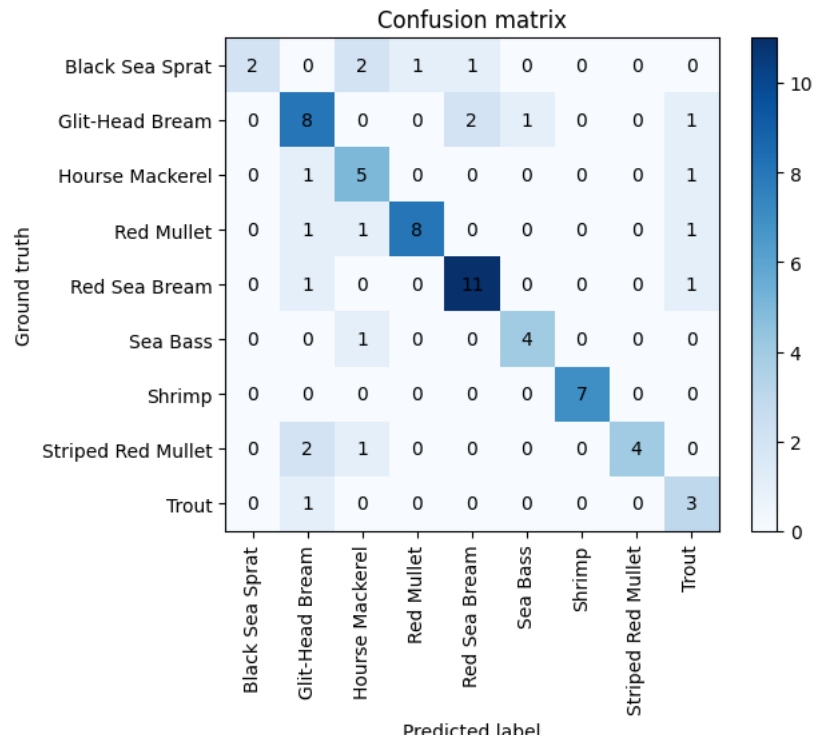
2/2 [=====] - 49s 6s/step
First 10 predicted value from model: [5 1 0 6 4 2 3 1 4 2]
First 10 ground truth label: [5, 1, 0, 6, 4, 2, 0, 4, 1, 3]

```

```

category = ['Black Sea Sprat', 'Glit-Head Bream', 'Hourse Mackerel', 'Red Mullet',
            'Red Sea Bream', 'Sea Bass', 'Shrimp', 'Striped Red Mullet', 'Trout']
conf_mat = confusion_matrix(y_true=test_batches.labels,
                            y_pred=np.argmax(y_hat, axis=1))
plt.imshow(conf_mat, cmap=plt.cm.Blues)
plt.colorbar()
indexes = np.arange(len(category))
for i in indexes:
    for j in indexes:
        plt.text(j, i, conf_mat[i, j],
                 horizontalalignment='center',
                 verticalalignment='center')
plt.xticks(indexes, category, rotation=90)
plt.xlabel('Predicted label')
plt.yticks(indexes, category)
plt.ylabel('Ground truth')
plt.title('Confusion matrix')
plt.show()

```



```
results = new_model.evaluate(test_batches, verbose=1)
print('Test accuracy: %.3f' % (results[1] * 100), '%')
print('Test loss: %.3f' % results[0])

2/2 [=====] - 46s 5s/step - loss: 0.9159 - accuracy: 0.7222
Test accuracy: 72.222 %
Test loss: 0.916
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.