

▼ Introduction

This kernel is a simple example of fruit classification with a simple CNN implemented with Keras model and tools. Click the blue "Edit Notebook" or "Fork Notebook" button at the top of this kernel to begin editing.

```
from google.colab import drive

drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive

▼ General Libs

```
# General Libs
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, BatchNormalization, |
from tensorflow.keras.optimizers import Adam
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
```

▼ Exploratory Analysis

Let's view some dataset samples.

This example uses keras ImageDataGenerator. This lib turns easier to read and use image classes from a subdirectoty structure.

▼ Reading the dataset

```
im_shape = (250,250)

TRAINING_DIR = '../content/drive/MyDrive/ds_frutas_am/train'
TEST_DIR = '../content/drive/MyDrive/ds_frutas_am/test'

seed = 10

BATCH_SIZE = 16
```

```
#Using keras ImageGenerator and flow_from_directory

# Subdivision in test/validation
data_generator = ImageDataGenerator(rescale=1./255, validation_split=0.2)
val_data_generator = ImageDataGenerator(rescale=1./255, validation_split=0.2)

# If you want data augmentation, uncomment and run this cell
data_generator = ImageDataGenerator(
    validation_split=0.2,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')

val_data_generator = ImageDataGenerator(rescale=1./255, validation_split=0.2)

# Generator para parte train
train_generator = data_generator.flow_from_directory(TRAINING_DIR, target_size=im_shape
                                                    class_mode='categorical', batch_si

# Generator para parte validação
validation_generator = val_data_generator.flow_from_directory(TRAINING_DIR, target_size
                                                            class_mode='categorical', batch_si

# Generator para dataset de teste
test_generator = ImageDataGenerator(rescale=1./255)
test_generator = test_generator.flow_from_directory(TEST_DIR, target_size=im_shape, shu
                                                    class_mode='categorical', batch_si

nb_train_samples = train_generator.samples
nb_validation_samples = validation_generator.samples
nb_test_samples = test_generator.samples
classes = list(train_generator.class_indices.keys())
print('Classes: '+str(classes))
num_classes = len(classes)

Found 72 images belonging to 6 classes.
Found 18 images belonging to 6 classes.
Found 30 images belonging to 6 classes.
Classes: ['acai', 'cupuacu', 'graviola', 'guarana', 'pupunha', 'tucuma']
```


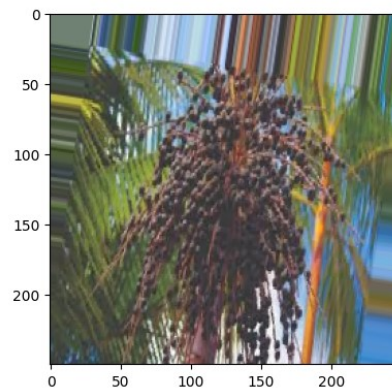
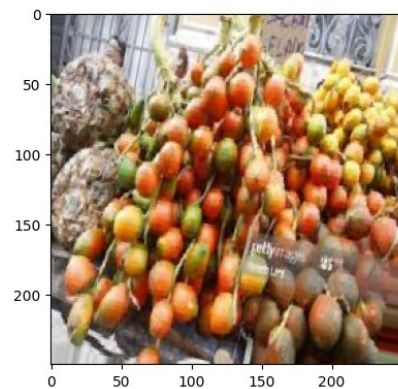


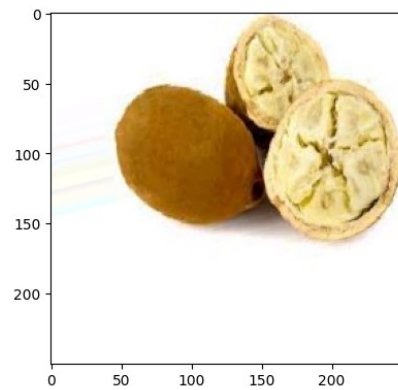
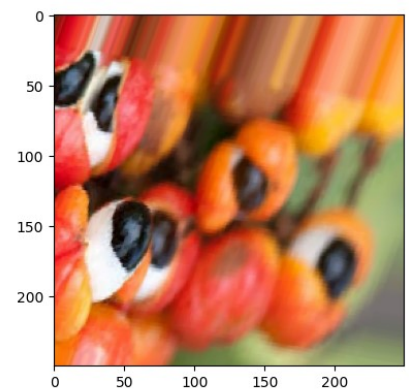

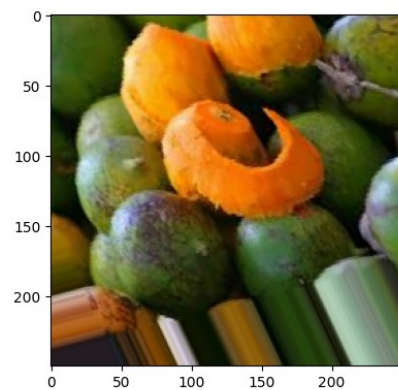
Showing some examples

```
# Visualizing some examples
plt.figure(figsize=(15,15))
for i in range(9):
```

```

#gera subfigures
plt.subplot(330 + 1 + i)
batch = train_generator.next()[0]*255
image = batch[0].astype('uint8')
plt.imshow(image)
plt.show()

```



Creating a simple CNN Model

```
model = Sequential()
model.add(Conv2D(20, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=(im_shape[0],im_shape[1],3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(40, kernel_size=(3,3), activation='relu'))
model.add(Flatten())
model.add(Dense(100, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
model.summary()

# Compila o modelo
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(),
              metrics=['accuracy'])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 248, 248, 20)	560
max_pooling2d (MaxPooling2D)	(None, 124, 124, 20)	0
conv2d_1 (Conv2D)	(None, 122, 122, 40)	7240
flatten (Flatten)	(None, 595360)	0
dense (Dense)	(None, 100)	59536100
dropout (Dropout)	(None, 100)	0
dense_1 (Dense)	(None, 6)	606
=====		
Total params: 59544506 (227.14 MB)		
Trainable params: 59544506 (227.14 MB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

epochs = 5

```
#Callback to save the best model
callbacks_list = [
    keras.callbacks.ModelCheckpoint(
        filepath='model.h5',
        monitor='val_loss', save_best_only=True, verbose=1),
    keras.callbacks.EarlyStopping(monitor='val_loss', patience=10, verbose=1)
]
```

... ..

```

# Training
history = model.fit(
    train_generator,
    steps_per_epoch=nb_train_samples // BATCH_SIZE,
    epochs=epochs,
    callbacks = callbacks_list,
    validation_data=validation_generator,
    verbose = 1,
    validation_steps=nb_validation_samples // BATCH_SIZE)

Epoch 1/5
4/4 [=====] - ETA: 0s - loss: 42.7706 - accuracy: 0.1250
Epoch 1: val_loss improved from inf to 30.65786, saving model to model.h5
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3000: UserWarning:
    saving_api.save_model(
4/4 [=====] - 13s 3s/step - loss: 42.7706 - accuracy: 0.1250
Epoch 2/5
4/4 [=====] - ETA: 0s - loss: 24.7431 - accuracy: 0.1607
Epoch 2: val_loss improved from 30.65786 to 4.14188, saving model to model.h5
4/4 [=====] - 13s 4s/step - loss: 24.7431 - accuracy: 0.1607
Epoch 3/5
4/4 [=====] - ETA: 0s - loss: 4.3579 - accuracy: 0.1406
Epoch 3: val_loss improved from 4.14188 to 1.59791, saving model to model.h5
4/4 [=====] - 15s 4s/step - loss: 4.3579 - accuracy: 0.1406
Epoch 4/5
4/4 [=====] - ETA: 0s - loss: 1.5971 - accuracy: 0.4107
Epoch 4: val_loss improved from 1.59791 to 1.55511, saving model to model.h5
4/4 [=====] - 12s 3s/step - loss: 1.5971 - accuracy: 0.4107
Epoch 5/5
4/4 [=====] - ETA: 0s - loss: 1.5341 - accuracy: 0.4107
Epoch 5: val_loss did not improve from 1.55511
4/4 [=====] - 9s 2s/step - loss: 1.5341 - accuracy: 0.4107

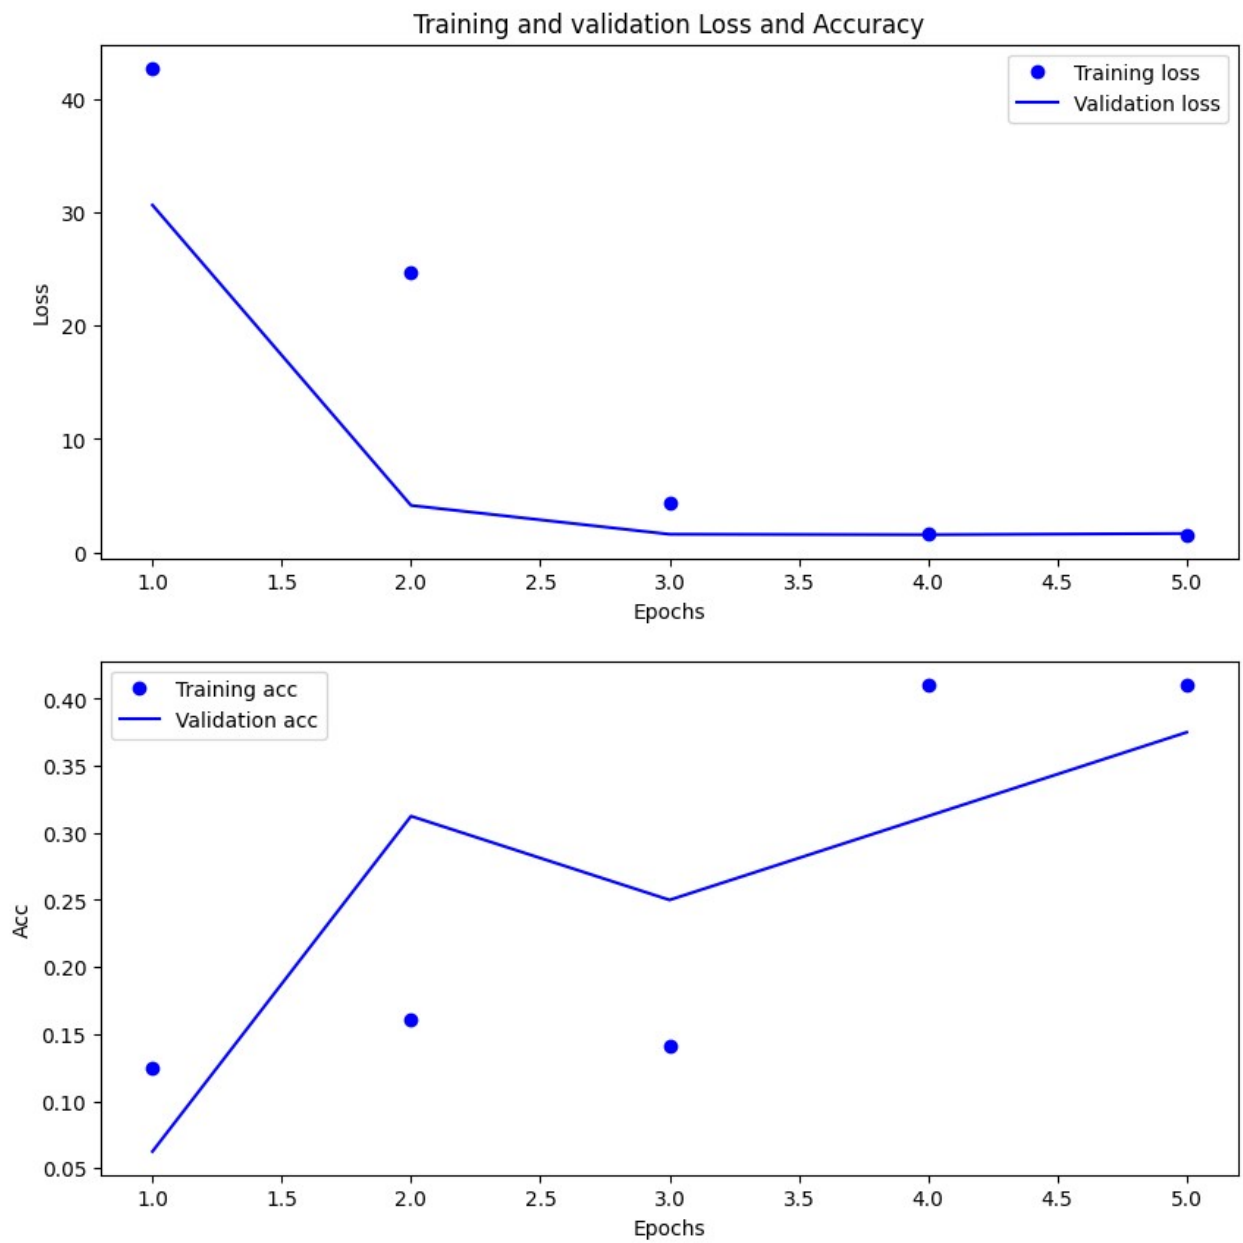
# Training curves
import matplotlib.pyplot as plt

history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs_x = range(1, len(loss_values) + 1)
plt.figure(figsize=(10,10))
plt.subplot(2,1,1)
plt.plot(epochs_x, loss_values, 'bo', label='Training loss')
plt.plot(epochs_x, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation Loss and Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.subplot(2,1,2)
acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']
plt.plot(epochs_x, acc_values, 'bo', label='Training acc')
plt.plot(epochs_x, val_acc_values, 'b', label='Validation acc')
# plt.title('Training and validation accuracy')
plt.xlabel('Epochs')

```

```
plt.xlabel('Epochs')  
plt.ylabel('Acc')  
plt.legend()  
plt.show()
```



Evaluating the model

```
# Load the best saved model
from tensorflow.keras.models import load_model

model = load_model('model.h5')

# Using the validation dataset
score = model.evaluate_generator(validation_generator)
print('Val loss:', score[0])
print('Val accuracy:', score[1])

<ipython-input-13-0b1386c018fa>:2: UserWarning: `Model.evaluate_generator` is deprecated
score = model.evaluate_generator(validation_generator)
Val loss: 1.57163667678833
Val accuracy: 0.2777777910232544

# Using the test dataset
score = model.evaluate_generator(test_generator)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

<ipython-input-14-8baddd65724c>:2: UserWarning: `Model.evaluate_generator` is deprecated
score = model.evaluate_generator(test_generator)
Test loss: 1.487316608428955
Test accuracy: 0.3333333432674408

import itertools

#Plot the confusion matrix. Set Normalize = True/False
def plot_confusion_matrix(cm, classes, normalize=True, title='Confusion matrix', cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.figure(figsize=(10,10))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        cm = np.around(cm, decimals=2)
        cm[np.isnan(cm)] = 0.0
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```
# Some reports
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

#On test dataset
Y_pred = model.predict_generator(test_generator)
y_pred = np.argmax(Y_pred, axis=1)
target_names = classes

#Confution Matrix
cm = confusion_matrix(test_generator.classes, y_pred)
plot_confusion_matrix(cm, target_names, normalize=False, title='Confusion Matrix')

#Classification Report
print('Classification Report')
print(classification_report(test_generator.classes, y_pred, target_names=target_names))
```

```
<ipython-input-16-4a63cf876f11>:6: UserWarning: `Model.predict_generator` is deprecated
Y_pred = model.predict_generator(test_generator)
```

```
Classification Report
```

	precision	recall	f1-score	support
acai	0.33	1.00	0.50	5
cupuacu	0.00	0.00	0.00	5
graviola	0.00	0.00	0.00	5
guarana	0.00	0.00	0.00	5
pupunha	0.36	1.00	0.53	5
tucuma	0.00	0.00	0.00	5
accuracy			0.33	30
macro avg	0.12	0.33	0.17	30
weighted avg	0.12	0.33	0.17	30

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: l
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: l
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: l
_warn_prf(average, modifier, msg_start, len(result))
```

