

```
!pip install -q git+https://github.com/tensorflow/docs

Preparing metadata (setup.py) ... done
Building wheel for tensorflow-docs (setup.py) ... done

!wget -q https://github.com/sayakpaul/Action-Recognition-in-TensorFlow/releases/download/v1.0.0/ucf101_top5.tar
!tar xf ucf101_top5.tar.gz
```

▼ Setup

```
from tensorflow_docs.vis import embed
from tensorflow import keras
from imutils import paths

import matplotlib.pyplot as plt
import tensorflow as tf
import pandas as pd
import numpy as np
import imageio
import cv2
import os
```

▼ Define hyperparameters

```
IMG_SIZE = 224
BATCH_SIZE = 64
EPOCHS = 10

MAX_SEQ_LENGTH = 20
NUM_FEATURES = 2048
```

▼ Data preparation

```
train_df = pd.read_csv("train.csv")
test_df = pd.read_csv("test.csv")

print(f"Total videos for training: {len(train_df)}")
print(f"Total videos for testing: {len(test_df)}")

train_df.sample(10)
```

```
Total videos for training: 594
Total videos for testing: 224
```

	video_name	tag	
459	v_ShavingBeard_g23_c03.avi	ShavingBeard	
135	v_PlayingCello_g10_c05.avi	PlayingCello	
75	v_CricketShot_g19_c02.avi	CricketShot	
189	v_PlayingCello_g18_c06.avi	PlayingCello	
411	v_ShavingBeard_g16_c03.avi	ShavingBeard	
37	v_CricketShot_g13_c03.avi	CricketShot	
52	v_CricketShot_g15_c04.avi	CricketShot	
101	v_CricketShot_g23_c03.avi	CricketShot	
395	v_ShavingBeard_g13_c06.avi	ShavingBeard	
489	v_TennisSwing_g09_c07.avi	TennisSwing	

```
# The following two methods are taken from this tutorial:
# https://www.tensorflow.org/hub/tutorials/action_recognition_with_tf_hub
```

```
def crop_center_square(frame):
    y, x = frame.shape[0:2]
    min_dim = min(y, x)
    start_x = (x // 2) - (min_dim // 2)
    start_y = (y // 2) - (min_dim // 2)
    return frame[start_y : start_y + min_dim, start_x : start_x + min_dim]
```

```
def load_video(path, max_frames=0, resize=(IMG_SIZE, IMG_SIZE)):
    cap = cv2.VideoCapture(path)
    frames = []
    try:
        while True:
            ret, frame = cap.read()
            if not ret:
                break
            frame = crop_center_square(frame)
            frame = cv2.resize(frame, resize)
            frame = frame[:, :, [2, 1, 0]]
            frames.append(frame)

            if len(frames) == max_frames:
                break
    finally:
        cap.release()
    return np.array(frames)
```

We can use a pre-trained network to extract meaningful features from the extracted frames. The [Keras Applications](#) module provides a number of state-of-the-art models pre-trained on the [ImageNet-1k dataset](#). We will be using the [InceptionV3 model](#) for this purpose.

```
def build_feature_extractor():
    feature_extractor = keras.applications.InceptionV3(
        weights="imagenet",
        include_top=False,
        pooling="avg",
        input_shape=(IMG_SIZE, IMG_SIZE, 3),
    )
    preprocess_input = keras.applications.inception_v3.preprocess_input

    inputs = keras.Input((IMG_SIZE, IMG_SIZE, 3))
    preprocessed = preprocess_input(inputs)

    outputs = feature_extractor(preprocessed)
    return keras.Model(inputs, outputs, name="feature_extractor")
```

```
feature_extractor = build_feature_extractor()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_orderin
87910968/87910968 [=====] - 1s 0us/step
```



The labels of the videos are strings. Neural networks do not understand string values, so they must be converted to some numerical form before they are fed to the model. Here we will use the [StringLookup](#) layer encode the class labels as integers.

```
label_processor = keras.layers.StringLookup(
    num_oov_indices=0, vocabulary=np.unique(train_df["tag"])
)
print(label_processor.get_vocabulary())

['CricketShot', 'PlayingCello', 'Punch', 'ShavingBeard', 'TennisSwing']
```

Finally, we can put all the pieces together to create our data processing utility.

```
def prepare_all_videos(df, root_dir):
    num_samples = len(df)
    video_paths = df["video_name"].values.tolist()
    labels = df["tag"].values
    labels = label_processor(labels[...], None).numpy()

    # `frame_masks` and `frame_features` are what we will feed to our sequence model.
    # `frame_masks` will contain a bunch of booleans denoting if a timestep is
    # masked with padding or not.
    frame_masks = np.zeros(shape=(num_samples, MAX_SEQ_LENGTH), dtype="bool")
    frame_features = np.zeros(
        shape=(num_samples, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32"
    )

    # For each video.
    for idx, path in enumerate(video_paths):
```

```
# Gather all its frames and add a batch dimension.
frames = load_video(os.path.join(root_dir, path))
frames = frames[None, ...]

# Initialize placeholders to store the masks and features of the current video.
temp_frame_mask = np.zeros(shape=(1, MAX_SEQ_LENGTH,), dtype="bool")
temp_frame_features = np.zeros(
    shape=(1, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32"
)

# Extract features from the frames of the current video.
for i, batch in enumerate(frames):
    video_length = batch.shape[0]
    length = min(MAX_SEQ_LENGTH, video_length)
    for j in range(length):
        temp_frame_features[i, j, :] = feature_extractor.predict(
            batch[None, j, :]
        )
    temp_frame_mask[i, :length] = 1 # 1 = not masked, 0 = masked

frame_features[idx,] = temp_frame_features.squeeze()
frame_masks[idx,] = temp_frame_mask.squeeze()

return (frame_features, frame_masks), labels

train_data, train_labels = prepare_all_videos(train_df, "train")
test_data, test_labels = prepare_all_videos(test_df, "test")

print(f"Frame features in train set: {train_data[0].shape}")
print(f"Frame masks in train set: {train_data[1].shape}")
```

```
1/1 [=====] - 0s 30ms/step
Frame features in train set: (594, 20, 2048)
Frame masks in train set: (594, 20)
```

The above code block will take ~20 minutes to execute depending on the machine it's being executed.

▼ The sequence model

Now, we can feed this data to a sequence model consisting of recurrent layers like GRU.

```
# Utility for our sequence model.
def get_sequence_model():
    class_vocab = label_processor.get_vocabulary()

    frame_features_input = keras.Input((MAX_SEQ_LENGTH, NUM_FEATURES))
    mask_input = keras.Input((MAX_SEQ_LENGTH,), dtype="bool")

    # Refer to the following tutorial to understand the significance of using `mask`:
    # https://keras.io/api/layers/recurrent_layers/gru/
    x = keras.layers.GRU(16, return_sequences=True)(
        frame_features_input, mask=mask_input
    )
    x = keras.layers.GRU(8)(x)
    x = keras.layers.Dropout(0.4)(x)
    x = keras.layers.Dense(8, activation="relu")(x)
    output = keras.layers.Dense(len(class_vocab), activation="softmax")(x)

    rnn_model = keras.Model([frame_features_input, mask_input], output)

    rnn_model.compile(
        loss="sparse_categorical_crossentropy", optimizer="adam", metrics=["accuracy"]
    )
    return rnn_model

# Utility for running experiments.
def run_experiment():
    filepath = "/tmp/video_classifier"
    checkpoint = keras.callbacks.ModelCheckpoint(
        filepath, save_weights_only=True, save_best_only=True, verbose=1
    )

    seq_model = get_sequence_model()
    history = seq_model.fit(
        [train_data[0], train_data[1]],
        train_labels,
        validation_split=0.3,
        epochs=EPOCHS,
        callbacks=[checkpoint],
    )

    seq_model.load_weights(filepath)
    _, accuracy = seq_model.evaluate([test_data[0], test_data[1]], test_labels)
    print(f"Test accuracy: {round(accuracy * 100, 2)}%")

    return history, seq_model

_, sequence_model = run_experiment()

Epoch 1/10
13/13 [=====] - ETA: 0s - loss: 1.2280 - accuracy: 0.4964
Epoch 1: val_loss improved from inf to 2.59972, saving model to /tmp/video_classifier
13/13 [=====] - 14s 232ms/step - loss: 1.2280 - accuracy: 0.4964 - val_loss: 2.5997 - val_accuracy: 0.2682
Epoch 2/10
13/13 [=====] - ETA: 0s - loss: 0.9668 - accuracy: 0.7229
Epoch 2: val_loss improved from 2.59972 to 2.40127, saving model to /tmp/video_classifier
13/13 [=====] - 0s 24ms/step - loss: 0.9668 - accuracy: 0.7229 - val_loss: 2.4013 - val_accuracy: 0.3464
Epoch 3/10
13/13 [=====] - ETA: 0s - loss: 0.8152 - accuracy: 0.8193
Epoch 3: val_loss did not improve from 2.40127
13/13 [=====] - 0s 21ms/step - loss: 0.8152 - accuracy: 0.8193 - val_loss: 2.7382 - val_accuracy: 0.3464
Epoch 4/10
12/13 [=====>...] - ETA: 0s - loss: 0.7344 - accuracy: 0.8776
Epoch 4: val_loss did not improve from 2.40127
13/13 [=====] - 0s 22ms/step - loss: 0.7298 - accuracy: 0.8795 - val_loss: 2.9071 - val_accuracy: 0.3464
Epoch 5/10
12/13 [=====>...] - ETA: 0s - loss: 0.6661 - accuracy: 0.8776
Epoch 5: val_loss did not improve from 2.40127
13/13 [=====] - 0s 32ms/step - loss: 0.6596 - accuracy: 0.8819 - val_loss: 3.0607 - val_accuracy: 0.3464
Epoch 6/10
11/13 [=====>....] - ETA: 0s - loss: 0.5712 - accuracy: 0.9290
```

```

Epoch 6: val_loss did not improve from 2.40127
13/13 [=====] - 0s 33ms/step - loss: 0.5633 - accuracy: 0.9277 - val_loss: 2.9756 - val_accuracy: 0.3464
Epoch 7/10
13/13 [=====] - ETA: 0s - loss: 0.5291 - accuracy: 0.9398
Epoch 7: val_loss did not improve from 2.40127
13/13 [=====] - 1s 40ms/step - loss: 0.5291 - accuracy: 0.9398 - val_loss: 3.0836 - val_accuracy: 0.3464
Epoch 8/10
11/13 [=====>....] - ETA: 0s - loss: 0.4780 - accuracy: 0.9545
Epoch 8: val_loss did not improve from 2.40127
13/13 [=====] - 1s 39ms/step - loss: 0.4737 - accuracy: 0.9566 - val_loss: 3.2253 - val_accuracy: 0.3464
Epoch 9/10
13/13 [=====] - ETA: 0s - loss: 0.4372 - accuracy: 0.9566
Epoch 9: val_loss did not improve from 2.40127
13/13 [=====] - 0s 33ms/step - loss: 0.4372 - accuracy: 0.9566 - val_loss: 3.2985 - val_accuracy: 0.3464
Epoch 10/10
13/13 [=====] - ETA: 0s - loss: 0.3902 - accuracy: 0.9398
Epoch 10: val_loss did not improve from 2.40127
13/13 [=====] - 0s 26ms/step - loss: 0.3902 - accuracy: 0.9398 - val_loss: 3.3919 - val_accuracy: 0.3464
7/7 [=====] - 0s 7ms/step - loss: 1.3408 - accuracy: 0.7455
Test accuracy: 74.55%

```

Note: To keep the runtime of this example relatively short, we just used a few training examples. This number of training examples is low with respect to the sequence model being used that has 99,909 trainable parameters. You are encouraged to sample more data from the UCF101 dataset using [the notebook](#) mentioned above and train the same model.

▼ Inference

```

def prepare_single_video(frames):
    frames = frames[None, ...]
    frame_mask = np.zeros(shape=(1, MAX_SEQ_LENGTH,), dtype="bool")
    frame_features = np.zeros(shape=(1, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32")

    for i, batch in enumerate(frames):
        video_length = batch.shape[0]
        length = min(MAX_SEQ_LENGTH, video_length)
        for j in range(length):
            frame_features[i, j, :] = feature_extractor.predict(batch[None, j, :])
            frame_mask[i, :length] = 1 # 1 = not masked, 0 = masked

    return frame_features, frame_mask

def sequence_prediction(path):
    class_vocab = label_processor.get_vocabulary()

    frames = load_video(os.path.join("test", path))
    frame_features, frame_mask = prepare_single_video(frames)
    probabilities = sequence_model.predict([frame_features, frame_mask])[0]

    for i in np.argsort(probabilities)[::-1]:
        print(f" {class_vocab[i]}: {probabilities[i] * 100:5.2f}%")
    return frames

# This utility is for visualization.
# Referenced from:
# https://www.tensorflow.org/hub/tutorials/action_recognition_with_tf_hub
def to_gif(images):
    converted_images = images.astype(np.uint8)
    imageio.mimsave("animation.gif", converted_images, duration=100)
    return embed.embed_file("animation.gif")

test_video = np.random.choice(test_df["video_name"].values.tolist())
print(f"Test video path: {test_video}")
test_frames = sequence_prediction(test_video)
to_gif(test_frames[:MAX_SEQ_LENGTH])

```

```
Test video path: v_ShavingBeard_g05_c02.avi
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 2s 2s/step
ShavingBeard: 51.49%
Punch: 16.89%
TennisSwing: 16.32%
CricketShot: 9.34%
PlayingCello: 5.96%
```

