

```

1 !pip install -q git+https://github.com/tensorflow/docs

    Preparing metadata (setup.py) ... done
    Building wheel for tensorflow-docs (setup.py) ... done

1 !wget -q https://github.com/sayakpaul/Action-Recognition-in-TensorFlow/releases/download/v1.0.0/ucf101_top5
2 !tar xf ucf101_top5.tar.gz

1 from tensorflow_docs.vis import embed
2 from tensorflow import keras
3 from imutils import paths
4
5 import matplotlib.pyplot as plt
6 import tensorflow as tf
7 import pandas as pd
8 import numpy as np
9 import imageio
10 import cv2
11 import os

1 IMG_SIZE = 224
2 BATCH_SIZE = 64
3 EPOCHS = 10
4
5 MAX_SEQ_LENGTH = 20
6 NUM_FEATURES = 2048

1 train_df = pd.read_csv("train.csv")
2 test_df = pd.read_csv("test.csv")
3 print(f"Total videos for training: {len(train_df)}")
4 print(f"Total videos for testing: {len(test_df)}")
5 train_df.sample(10)

    Total videos for training: 594
    Total videos for testing: 224

```

	video_name	tag
448	v_ShavingBeard_g21_c06.avi	ShavingBeard
525	v_TennisSwing_g15_c01.avi	TennisSwing
192	v_PlayingCello_g19_c02.avi	PlayingCello
581	v_TennisSwing_g23_c06.avi	TennisSwing
588	v_TennisSwing_g25_c01.avi	TennisSwing
407	v_ShavingBeard_g15_c06.avi	ShavingBeard
359	v_ShavingBeard_g08_c01.avi	ShavingBeard
207	v_PlayingCello_g21_c03.avi	PlayingCello
222	v_PlayingCello_g23_c06.avi	PlayingCello
585	v_TennisSwing_g24_c04.avi	TennisSwing

```

1 def crop_center_square(frame):
2     y, x = frame.shape[0:2]
3     min_dim = min(y, x)
4     start_x = (x // 2) - (min_dim // 2)
5     start_y = (y // 2) - (min_dim // 2)
6     return frame[start_y : start_y + min_dim, start_x : start_x + min_dim]
7 def load_video(path, max_frames=0, resize=(IMG_SIZE, IMG_SIZE)):
8     cap = cv2.VideoCapture(path)
9     frames = []
10    try:
11        while True:
12            ret, frame = cap.read()
13            if not ret:
14                break
15            frame = crop_center_square(frame)
16            frame = cv2.resize(frame, resize)
17            frame = frame[:, :, [2, 1, 0]]
18            frames.append(frame)

```

```

19
20         if len(frames) == max_frames:
21             break
22     finally:
23         cap.release()
24     return np.array(frames)
25
26
27 def build_feature_extractor():
28     feature_extractor = keras.applications.InceptionV3(
29         weights="imagenet",
30         include_top=False,
31         pooling="avg",
32         input_shape=(IMG_SIZE, IMG_SIZE, 3),
33     )
34     preprocess_input = keras.applications.inception_v3.preprocess_input
35
36     inputs = keras.Input((IMG_SIZE, IMG_SIZE, 3))
37     preprocessed = preprocess_input(inputs)
38
39     outputs = feature_extractor(preprocessed)
40     return keras.Model(inputs, outputs, name="feature_extractor")
41
42 feature_extractor = build_feature_extractor()

```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception\_v3/inception\_v3\_weights\_tf\_dim\_ordering\_tf\_87910968/87910968 [=====] - 0s 0us/step
```



```
1 label_processor = keras.layers.StringLookup(
2     num_oov_indices=0, vocabulary=np.unique(train_df["tag"])
3 )
4 print(label_processor.get_vocabulary())

['CricketShot', 'PlayingCello', 'Punch', 'ShavingBeard', 'TennisSwing']
```

```

1 def prepare_all_videos(df, root_dir):
2     num_samples = len(df)
3     video_paths = df["video_name"].values.tolist()
4     labels = df["tag"].values
5     labels = label_processor(labels[...], None).numpy()
6     frame_masks = np.zeros(shape=(num_samples, MAX_SEQ_LENGTH), dtype="bool")
7     frame_features = np.zeros(
8         shape=(num_samples, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32"
9     )
10    for idx, path in enumerate(video_paths):
11        frames = load_video(os.path.join(root_dir, path))
12        frames = frames[None, ...]
13        temp_frame_mask = np.zeros(shape=(1, MAX_SEQ_LENGTH,), dtype="bool")
14        temp_frame_features = np.zeros(
15            shape=(1, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32"
16        )
17        for i, batch in enumerate(frames):
18            video_length = batch.shape[0]
19            length = min(MAX_SEQ_LENGTH, video_length)
20            for j in range(length):
21                temp_frame_features[i, j, :] = feature_extractor.predict(
22                    batch[None, j, :])
23            )
24            temp_frame_mask[i, :length] = 1
25
26        frame_features[idx,] = temp_frame_features.squeeze()
27        frame_masks[idx,] = temp_frame_mask.squeeze()
28
29    return (frame_features, frame_masks), labels
30 train_data, train_labels = prepare_all_videos(train_df, "train")
31 test_data, test_labels = prepare_all_videos(test_df, "test")
32 print(f"Frame features in train set: {train_data[0].shape}")
33 print(f"Frame masks in train set: {train_data[1].shape}")

```



```

1/1 [=====] - 0s 73ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 77ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 64ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step

```

```

1 def get_sequence_model():
2     class_vocab = label_processor.get_vocabulary()
3     frame_features_input = keras.Input((MAX_SEQ_LENGTH, NUM_FEATURES))
4     mask_input = keras.Input((MAX_SEQ_LENGTH,), dtype="bool")
5     x = keras.layers.GRU(16, return_sequences=True)(
6         frame_features_input, mask=mask_input
7     )
8     x = keras.layers.GRU(8)(x)
9     x = keras.layers.Dropout(0.4)(x)
10    x = keras.layers.Dense(8, activation="relu")(x)
11    output = keras.layers.Dense(len(class_vocab), activation="softmax")(x)
12    rnn_model = keras.Model([frame_features_input, mask_input], output)
13    rnn_model.compile(
14        loss="sparse_categorical_crossentropy", optimizer="adam", metrics=["accuracy"]
15    )
16    return rnn_model
17 def run_experiment():
18     filepath = "/tmp/video_classifier"
19     checkpoint = keras.callbacks.ModelCheckpoint(
20         filepath, save_weights_only=True, save_best_only=True, verbose=1
21     )
22
23     seq_model = get_sequence_model()
24     history = seq_model.fit(
25         [train_data[0], train_data[1]],
26         train_labels,
27         validation_split=0.3,

```

```

28     epochs=EPOCHS,
29     callbacks=[checkpoint],
30 )
31
32 seq_model.load_weights(filepath)
33 _, accuracy = seq_model.evaluate([test_data[0], test_data[1]], test_labels)
34 print(f"Test accuracy: {round(accuracy * 100, 2)}%")
35
36 return history, seq_model
37
38
39 _, sequence_model = run_experiment()

Epoch 1/10
12/13 [=====>...] - ETA: 0s - loss: 1.3072 - accuracy: 0.4844
Epoch 1: val_loss improved from inf to 1.63220, saving model to /tmp/video_classifier
13/13 [=====] - 14s 271ms/step - loss: 1.2843 - accuracy: 0.5012 - val_loss: 1.6322 - val_accuracy: 0.1397
Epoch 2/10
13/13 [=====] - ETA: 0s - loss: 0.9665 - accuracy: 0.7181
Epoch 2: val_loss improved from 1.63220 to 1.58667, saving model to /tmp/video_classifier
13/13 [=====] - 0s 38ms/step - loss: 0.9665 - accuracy: 0.7181 - val_loss: 1.5867 - val_accuracy: 0.2011
Epoch 3/10
12/13 [=====>...] - ETA: 0s - loss: 0.8751 - accuracy: 0.8073
Epoch 3: val_loss did not improve from 1.58667
13/13 [=====] - 0s 34ms/step - loss: 0.8650 - accuracy: 0.8193 - val_loss: 1.8510 - val_accuracy: 0.2514
Epoch 4/10
12/13 [=====>...] - ETA: 0s - loss: 0.7857 - accuracy: 0.8333
Epoch 4: val_loss did not improve from 1.58667
13/13 [=====] - 0s 33ms/step - loss: 0.7725 - accuracy: 0.8386 - val_loss: 1.8268 - val_accuracy: 0.2905
Epoch 5/10
13/13 [=====] - ETA: 0s - loss: 0.7205 - accuracy: 0.8723
Epoch 5: val_loss did not improve from 1.58667
13/13 [=====] - 0s 22ms/step - loss: 0.7205 - accuracy: 0.8723 - val_loss: 1.9336 - val_accuracy: 0.3017
Epoch 6/10
13/13 [=====] - ETA: 0s - loss: 0.6462 - accuracy: 0.9229
Epoch 6: val_loss did not improve from 1.58667
13/13 [=====] - 0s 21ms/step - loss: 0.6462 - accuracy: 0.9229 - val_loss: 1.9797 - val_accuracy: 0.3240
Epoch 7/10
13/13 [=====] - ETA: 0s - loss: 0.5791 - accuracy: 0.9253
Epoch 7: val_loss did not improve from 1.58667
13/13 [=====] - 0s 21ms/step - loss: 0.5791 - accuracy: 0.9253 - val_loss: 2.0976 - val_accuracy: 0.3128
Epoch 8/10
13/13 [=====] - ETA: 0s - loss: 0.5171 - accuracy: 0.9542
Epoch 8: val_loss did not improve from 1.58667
13/13 [=====] - 0s 21ms/step - loss: 0.5171 - accuracy: 0.9542 - val_loss: 2.1763 - val_accuracy: 0.3296
Epoch 9/10
13/13 [=====] - ETA: 0s - loss: 0.4852 - accuracy: 0.9398
Epoch 9: val_loss did not improve from 1.58667
13/13 [=====] - 0s 21ms/step - loss: 0.4852 - accuracy: 0.9398 - val_loss: 2.2442 - val_accuracy: 0.3128
Epoch 10/10
13/13 [=====] - ETA: 0s - loss: 0.4472 - accuracy: 0.9566
Epoch 10: val_loss did not improve from 1.58667
13/13 [=====] - 0s 21ms/step - loss: 0.4472 - accuracy: 0.9566 - val_loss: 2.2161 - val_accuracy: 0.3184
7/7 [=====] - 0s 7ms/step - loss: 1.0930 - accuracy: 0.7054
Test accuracy: 70.54%

```

```

1 def prepare_single_video(frames):
2     frames = frames[None, ...]
3     frame_mask = np.zeros(shape=(1, MAX_SEQ_LENGTH,), dtype="bool")
4     frame_features = np.zeros(shape=(1, MAX_SEQ_LENGTH, NUM_FEATURES), dtype="float32")
5
6     for i, batch in enumerate(frames):
7         video_length = batch.shape[0]
8         length = min(MAX_SEQ_LENGTH, video_length)
9         for j in range(length):
10             frame_features[i, j, :] = feature_extractor.predict(batch[None, j, :])
11             frame_mask[i, :length] = 1 # 1 = not masked, 0 = masked
12
13     return frame_features, frame_mask
14 def sequence_prediction(path):
15     class_vocab = label_processor.get_vocabulary()
16
17     frames = load_video(os.path.join("test", path))
18     frame_features, frame_mask = prepare_single_video(frames)
19     probabilities = sequence_model.predict([frame_features, frame_mask])[0]
20
21     for i in np.argsort(probabilities)[::-1]:
22         print(f" {class_vocab[i]}: {probabilities[i] * 100:5.2f}%")
23     return frames
24 def to_gif(images):
25     animated_images = images.astype(np.uint8)

```

```

25     converted_images = images.astype(np.uint8)
26     imageio.mimsave("animation.gif", converted_images, duration=100)
27     return embed.embed_file("animation.gif")
28
29
30 test_video = np.random.choice(test_df["video_name"].values.tolist())
31 print(f"Test video path: {test_video}")
32 test_frames = sequence_prediction(test_video)
33 to_gif(test_frames[:MAX_SEQ_LENGTH])

```

Test video path: v_ShavingBeard_g04_c01.avi

```

1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 2s 2s/step

```

ShavingBeard: 26.35%

PlayingCello: 26.27%

Punch: 18.13%

CricketShot: 15.62%

TennisSwing: 13.62%

